

November 20, 2019  
DRAFT

# **Byzantine Agreement with Incomplete Views**

Hanjun Li

November 20, 2019

School of Computer Science  
Carnegie Mellon University  
Pittsburgh, PA 15213

**Thesis Committee:**

Vipul Goyal  
Bryan Parno

*Submitted in partial fulfillment of the requirements  
for the degree of Master of Science.*

November 20, 2019  
DRAFT

**Keywords:** Byzantine Agreement

November 20, 2019  
DRAFT

*For my dog*

## **Abstract**

A short summary

## **Acknowledgments**

My advisor, and my collaborator

November 20, 2019  
DRAFT

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	Related Works . . . . .	1
1.2	Contribution . . . . .	1
<b>2</b>	<b>Preliminaries</b>	<b>3</b>
2.1	Notations . . . . .	3
2.2	Unique Digital Signature and PKI . . . . .	3
2.3	Verifiable Random Function . . . . .	4
2.4	Our Model and Main Results . . . . .	6
<b>3</b>	<b>Positive Result</b>	<b>9</b>
3.1	Graded Broadcast . . . . .	9
3.2	Oblivious Leader Selection . . . . .	11
3.3	Main protocol . . . . .	13
<b>4</b>	<b>Negative Result</b>	<b>15</b>
4.1	Broadcast from Byzantine Agreement . . . . .	15
4.2	Negative Result for Broadcast . . . . .	15
<b>5</b>	<b>Extension: Unidirectional Edges</b>	<b>17</b>
	<b>Bibliography</b>	<b>20</b>

# Chapter 1

## Introduction

### 1.1 Related Works

**Introduced; deterministic;  $1/3$**  [LSP], [PSL80], [FL82], [GM98],

**Authenticated;  $1/2$**  [LSP], [PSL80], [DS83],

**Randomized;  $1/3$  and  $1/2$**  [BO83], [Rab83], [FM97], [KK06], [Mic17],

**Imcomplete Network** [KS09], [CGO10], [Dol82],

### 1.2 Contribution



November 20, 2019  
DRAFT

# Chapter 2

## Preliminaries

### 2.1 Notations

We use  $\{0, 1\}^*$  to denote the set of finite binary strings. When describing a probabilistic algorithm  $F(\cdot)$ ,  $F(x)$  refers to the probability space that assigns any string  $y$  the probability that  $F$  on input  $x$  outputs  $y$ . We write  $y \xleftarrow{R} F(x)$  to describe assigning to  $y$  an element randomly selected according to  $F(x)$ . In contrast, for a deterministic algorithm  $G(\cdot)$ , we simply write  $y = G(x)$  to describe the output of  $G$  on input  $x$  being  $y$ .

We will use the usual notation  $G = (V, E)$  to represent the communication graph among the participants.  $V$  is the set of nodes in the graph. Each node  $P_i$  represents the participant  $P_i$ .  $E$  is the set of edges. Each edge  $(P_j, P_k)$  represents the fact that  $P_k$  is in the *view* of  $P_j$  (see Section 2.2). Our model assumes only bi-directional edges. That is, if  $P_k$  is in the view of  $P_j$  then  $P_j$  is also in the view of  $P_k$ . We write  $\Gamma_i$  as the inclusive neighbor of  $P_i$  in  $G$  (i.e. all participants in the view of  $P_i$  including itself).

### 2.2 Unique Digital Signature and PKI

Similar to other authenticated Byzantine Agreement protocols, we will use a Digital Signature scheme to ensure that a corrupted participant can either choose to forward or ignore a signed message, but can never forge a signed message. In addition, we also require that for any public key and any message  $m$ , there is a unique valid signature. The additional uniqueness constraint is introduced in [GO92], and a construction based on the RSA assumption is provided in [MRVil]. Briefly, we summarize the notion of Unique Digital Signature below.

**Notation:** A Unique Digital Signature scheme is a triple of polynomial time computable algorithms  $(Gen, Sign, Verify)$  as described below.

- $Gen(\cdot)$  is a probabilistic algorithm that takes a unary string of length  $k$ , the security parameter, as input, and outputs two binary strings, a public key  $P_k$  and a secret key  $S_k$ .

We write this as  $(S_k, P_k) \xleftarrow{R} Gen(1^k)$ .

- $Sign(\cdot, \cdot)$  is a deterministic algorithm that takes in the secret key  $S_k$  and a message  $x$ , and produces a signature  $Sig_{P_k}(x)$ .  
We write this as  $Sig = Sign(S_k, x)$ .
- $Verify(\cdot, \cdot, \cdot)$  is a probabilistic algorithm that takes a public key  $P_k$ , a message  $x$  and a signature  $Sig$  as input, and outputs a bit  $b \in \{0, 1\}$ .  
We write this as  $b \stackrel{R}{\leftarrow} Verify(P_k, x, Sig)$ .

We now briefly describe the correctness and the security of a Unique Digital Signature scheme.

**Correctness:**

- $Verify$  accepts (i.e. outputs 1) a valid signature produced by  $Sign$  with overwhelming probability over  $(P_k, S_k) \stackrel{R}{\leftarrow} Gen(1^k)$
- There do not exist values  $(P_k, x, Sig_1, Sig_2)$  such that  $Sig_1 \neq Sig_2$ , and  $Verify(P_k, x, Sig_1) = Verify(P_k, x, Sig_2) = 1$ .

**Security:** Let  $T$  be any polynomial time algorithm. The probability that  $T$  wins the following game must be negligible:

- Run  $(P_k, S_k) \stackrel{R}{\leftarrow} Gen(1^k)$ :
- Run  $(x, Sig) \stackrel{R}{\leftarrow} T^{Sign(S_k, \cdot)}(1^k, P_k)$
- $T$  wins if  $1 \stackrel{R}{\leftarrow} Verify(P_k, x, Sig)$

This security definition is also called existentially unforgeable.

**Public-key Infrastructure:** For a Digital Signature scheme to be useful for authenticating messages, a trusted setup phase is required to first assign each participant  $P_i$  its key pair  $(P_{k_i}, S_{k_i})$ , and next distribute public keys of the participants. This setup phase is generally referred to as a Public-key infrastructure (PKI) setup. In our relaxed model, the setup phase only assigns to each participants a subset of the public keys. If the public key of  $P_j$  is assigned to  $P_i$ , we say  $P_i$  *trusts*  $P_j$ . The set of *trusted* participants by  $P_i$  is called the *view* of  $P_i$ . In the later sections, we will use the player id to identify its public key. For example a signature by  $P_i$  on some message  $m_i$  will be written as  $Sig_i(m_i)$ .

## 2.3 Verifiable Random Function

When describing and analyzing our protocols, we will assume that every participant has access to a public random function  $H$ , mapping  $\{0, 1\}^*$  to  $\{0, 1\}^k$  for any  $k$ . In this idealized model, when  $H$  receives a query string  $x$ , it selects a  $k$  bit string uniformly at random as its output  $H(x)$ . All further query of  $x$  all result in the same output  $H(x)$ .

Note that the random oracle  $H$  is only introduced to simplify our analysis. The usage of  $H$  can be replaced by a verifiable random function (VRF) scheme which can be constructed under the RSA assumption [MRVil]. We now briefly summarize the notion of VRF below.

**Notation:** A Verifiable Random Function scheme is a triple of polynomial time computable algorithms  $(G, F, V)$  as described below.

- $G(\cdot)$  is a probabilistic algorithm that takes a unary string of length  $k$ , the security parameter, as input, and outputs two binary strings, a public key  $PK$  and a secret key  $SK$ .

We write this as  $(PK, SK) \xleftarrow{R} G(1^k)$

- $F(\cdot, \cdot)$  is a deterministic algorithm that takes two binary strings, the secret key  $SK$  and a seed  $x$ , and outputs two binary strings, the value  $v$  and its corresponding proof  $proof$ .

We write this as  $(v, proof) = F(SK, x)$ . For convenience, we sometimes write  $F = (F_1, F_2)$  where  $v = F_1(SK, x)$  and  $proof = F_2(SK, x)$ .

- $V(\cdot, \cdot, \cdot, \cdot)$  is a probabilistic algorithm that takes four binary strings, the public key  $PK$ , the seed  $x$ , the value  $v$ , and the proof  $proof$ , as input, and output a bit  $b \in \{1, 0\}$ .

We write this as  $b \xleftarrow{R} V(PK, x, v, proof)$

At a high level, we can think of  $G$  as the function generator,  $F$  as the function evaluator, and  $V$  as the function verifier. We now describe the correctness and the security of a VRF scheme.

**Correctness:** The following must hold with overwhelming probability over  $(PK, SK) \xleftarrow{R} G(1^k)$ :

- For all  $x$  in its domain,  $F_1(SK, x)$  produces a string in its correct range.
- For all  $x$  in its domain, if  $(v, proof) = F(SK, x)$ , then  $1 \xleftarrow{R} V(PK, x, v, proof)$ .
- For every  $x, v_1, v_2, proof_1, proof_2$  such that  $v_1 \neq v_2$ , either  $0 \xleftarrow{R} V(PK, x, v_1, proof_1)$  or  $0 \xleftarrow{R} V(PK, x, v_2, proof_2)$ .

At a high level, the correctness requires that the proof produced by  $F_2$  can be uniquely verified by  $V$ .

**Security:** Let  $T = (T_E, T_J)$  be any pair of polynomial time algorithm (in the security parameter  $k$ ). The advantage that  $T$  has in succeeding the following game over  $1/2$  (i.e. random guessing) must be negligible.

- Run  $(PK, SK) \xleftarrow{R} G(1^k)$
- Run  $(x, state) \xleftarrow{R} T_E^{F(SK, \cdot)}(1^k, PK)$
- Randomly choose a bit  $r \xleftarrow{R} \{0, 1\}$ :
  - if  $r = 0$ , run  $v = F_1(SK, x)$
  - if  $r = 1$ , sample  $v$  at random from the range of  $F_1$ .
- Run  $b \xleftarrow{R} T_J^{F(SK, \cdot)}(1^k, v, state)$  where  $b$  is the guess of  $T$ .  $T$  succeeds if  $x$  is in the domain of  $F_1$ ,  $x$  is not asked as a query to  $F(SK, \cdot)$  by  $T_E$  or  $T_J$ , and  $b = r$ .

At a high level, the security requires that the output of  $F_1$  is indistinguishable from a random string.

**Using VRF in place of  $H(\cdot)$ :** We now describe how to use a VRF scheme to simulate the random oracle  $H$  used in our protocol. In the trusted PKI setup phase, each participant  $P_i$  is assigned its own key pair  $(PK_i, SK_i)$ , and also the set of public keys in its view. In our protocol, whenever a participant  $P_i$  receives a value  $v_j$  originated from  $P_j$ , and needs to evaluate  $H(v_j)$ , we instead ask  $P_j$  to attach the evaluation  $(r_j, proof_j) = F(SK_j, v_j)$  and its public key  $PK_j$  to  $v_j$ . That is,  $P_i$  will actually receive  $v'_j = (v_j, r_j, proof_j, PK_j)$ , and can simply run  $V(PK_j, v_j, r_j, proof_j)$  to verify that  $r_j$  is the correct result. Each message is additionally signed by  $P_j$  so that no one can forge an evaluation of  $H(v_j)$ .

In the case where  $P_j$  is not trusted by  $P_i$ , (hence all messages of  $v_j$  is forwarded by some other participant) our protocol always guarantees that at least one forwarding is from an honest participant, who has verified its signature and the included public key  $PK_i$ . Whenever  $P_j$  receives two contradicting  $PK_i$  and  $PK'_i$ ,  $P_j$  discards all messages from  $P_i$  since it must have been corrupted.

## 2.4 Our Model and Main Results

Our model is a relaxation from the standard one assumed in most authenticated Byzantine Agreement protocols (e.g. [Mic17], [KK06]). Details are provided below.

**Communication:** All participants communicate in synchronized rounds, over authenticated and private point-to-point channels. Note that such authenticated channels only exists between two parties who trusts each other, so our communication network is not a *complete* network. Messages are sent as the start of a round, and received by the end of the round.

**The Adversary:** The adversary is a polynomial time algorithm, that can adaptively corrupt honest participants during the protocol. Corrupted participants can deviate from the protocol in arbitrary ways. At the start of any round, the adversary may corrupt additional players before receiving messages from all honest participants, and then decide what to send from all corrupted participants. The adversary knows the public keys from all participants at the beginning.

**Results:** We first state the standard definition of a Byzantine Agreement protocol. For simplicity, we only consider the *binary* version.

**Definition 1.** (*Byzantine Agreement*) For a set of participants  $P_1, \dots, P_N$ , where each  $P_i$  holds an initial input  $v_i \in \{0, 1\}$ , when the protocol terminates, the following conditions must hold for any adversary:

- (*Validity*) If all honest participants begin with the same input  $v$ , they also output  $v$ .
- (*Agreement*) All honest participants output the same value

Let  $\alpha_i$  be the fraction of corrupted participants in the view of an honest participant  $P_i$ . We define  $\alpha \equiv \max_i \{\alpha_i\}$ . Let  $\delta_{ij}$  be the fraction of overlapping between the views of two honest participants  $P_i, P_j$  (i.e.  $\delta_{ij} = |\Gamma_i \cap \Gamma_j|/|\Gamma_i|$ ). We similarly define  $\delta \equiv \min_{i,j} \{\delta_{ij}\}$ . Now we are ready to state our results as the following theorems:

**Theorem 1.** *If  $\delta > 2\alpha$ , there exists an expected  $O(n)$  round Byzantine Agreement protocol. Further if  $\alpha = 1/2 - \epsilon$  for any constant  $\epsilon$ , there exists an expected constant round Byzantine Agreement protocol in our model.*

**Theorem 2.** *If  $\alpha \geq 1/2$  or  $\delta \leq 2\alpha$ , there does not exist a Byzantine Agreement protocol in our model*

November 20, 2019  
DRAFT

## Chapter 3

### Positive Result

This section first introduces several subprotocols as building blocks, and in the end use them in our main protocol. For simplicity, in the following discussion we will assume that all honest participants have a view of fixed size  $n$ . However, we note that all of our results hold without this restriction. Intuitively, since every pair of honest participants has at least a  $\delta$  overlapping in their views, the size of their views can not differ by too much.

#### 3.1 Graded Broadcast

A Graded Broadcast protocol is a similar but weaker notion to broadcast. At a high level, it simulates a broadcast with some participants fail to receive the message. It, however, guarantees that all participants that successfully receive a message indeed receive the same message. In addition, if the sender is an honest participant, the broadcast always succeeds. The name “Graded Broadcast” comes from the way a participant decide whether to accept a message: the successful message is assigned a positive grade, while a failed message is assigned grade 0. This idea is first introduced in [FM97] in the standard model. We now give a formal definition in our model.

**Definition 2.** (*Graded Broadcast*) For a set of participants  $S = \{P_1, \dots, P_N\}$ , and a distinguished dealer  $P_d \in S$  holding an initial message  $m$ , when the protocol terminates, the following conditions must hold for any adversary:

- Each honest participant  $P_i$  in the view of  $P_d$  (i.e.  $P_i \in \Gamma_d$ ) outputs  $(m_i, g_i)$ , where  $g_i \in \{0, 1\}$ .
- (Validity) If  $P_d$  is honest, then  $m_i = m$ , and  $g_i = 1$  for all  $P_i \in \Gamma_d$ .
- (Agreement) If two honest participants  $P_i, P_j \in \Gamma_d$  outputs  $(m_i, 1)$ , and  $(m_j, 1)$ , then  $m_i = m_j$ .

Next we give our protocol (Algorithm 1) that achieves Graded Broadcast assuming  $\delta \geq 2\alpha$ .



---

**Algorithm 1** Graded Broadcast

---

- 1: The dealer  $P_d$  signs message  $m$ , and sends  $(m, \text{Sig}_d(m))$  to all participants in its view.
  - 2: Every honest participant  $P_i$  in the view of  $P_d$  receives  $(m, \text{Sig}_d(m))$ , and verifies the signature.
    - if the signature is valid,  $P_i$  forwards it to all participants in the view of  $P_i$ .
    - else  $P_i$  does nothing.
  - 3: Every honest participant  $P_i$  in the view of  $P_d$  receives forwarded messages  $m_j, \text{Sig}_d(m_j)$  from some  $P_j$  in its own view.
    - If there are contradicting valid signatures of  $m_j \neq m_k$  for some  $j, k$ , then  $P_i$  outputs  $(\phi, 0)$ .
    - If there are less than  $(\delta - \alpha)n$  number of valid signatures on some message  $m$ , then outputs  $(\phi, 0)$ .
    - Else, outputs  $(m, 1)$ .
- 

We now show the following claims about Algorithm 1.

**Claim 3.** (Validity) *If the dealer  $P_d$  is honest with message  $m$ , then all honest participants  $P_i$  in the view of  $P_d$  output  $(m, 1)$  at the end of the Graded Broadcast protocol in Algorithm 1.*

*Proof.* Since  $P_d$  is honest, only signatures of the message  $m$  is ever sent out by  $P_d$ . By the security of a Digital Signature scheme, no contradicting signature (on a different message  $m' \neq m$ ) can be forged. Therefore, the first “If” branch in **Step 3** will not happen.

By our model assumption, there are at least  $(1 - \alpha)n$  honest participants in the view of  $P_d$ . For any honest  $P_i$ , at least  $((1 - \alpha) - (1 - \delta))n = (\delta - \alpha)n$  of them are also in the view of  $P_i$ . Therefore, any honest participant  $P_i$  in the view of  $P_d$  will receive at least  $(\delta - \alpha)n$  forwarded messages with valid signatures on  $m$ . Therefore, the second “If” branch in **Step 3** will also not happen, and  $P_i$  outputs  $(m, 1)$ .  $\square$

**Claim 4.** (Agreement) *If two honest participants  $P_i, P_j$  in the view of  $P_d$  output  $(m_i, 1)$  and  $(m_j, 1)$ , and if  $\delta > 2\alpha$ , then  $m_i = m_j$ .*

*Proof.* Suppose  $m_i \neq m_j$ . In **Step 3**,  $P_i$  has received at least  $(\delta - \alpha)n$  valid signatures on  $m_i$ . By assumption,  $(\delta - \alpha)n > \alpha n$ , so at least one of them is forwarded by some honest participant  $P_k$  in the view of  $P_d$ .

Since  $P_k$  is honest, in **Step 2**,  $P_k$  must have also sent  $(m_i, \text{Sig}_d(m_i))$  to  $P_j$ , and  $P_j$  must have seen contradicting valid signatures of  $m_i \neq m_j$  in **Step 3**. This leads to a contradiction.  $\square$

The above claims lead to the following lemma that we will use as a building block to prove Theorem 1.

**Lemma 5.** *If  $\delta > 2\alpha$ , there exists a three round Graded Broadcast protocol.*

*Proof.* By Algorithm 1, Claim 3, and Claim 4.  $\square$

## 3.2 Oblivious Leader Selection

An Oblivious Leader Selection protocol is used for electing an honest leader that is agreed on by all honest participants. However, we only need this to happen with *some* probability, to which we refer as the *fairness* of the protocol. If the fairness is a constant, then running the protocol repeatedly will give us an honest leader in expected constant round. If the fairness is  $\Omega(1/n)$ , then running the protocol repeatedly will give us an honest leader in expected  $O(n)$  round. Our definition is similar to the one present in [KK06]. Our protocol construction is inspired by the *ConcreteCoin* protocol present in [Mic17]. We now give a formal definition in our model:

**Definition 3.** (*Oblivious Leader Selection*) For a set of participants  $P_1, \dots, P_N$  and fairness  $\gamma$ , when the protocol terminates, the following conditions must hold with probability at least  $\gamma$ :

- Every honest participant  $P_i$  outputs  $P_l$  and  $P_l$  is honest by the end of the protocol.

When such an event happens, we say that an honest leader  $P_l$  is elected.

Next, we give our protocol (Algorithm 2) that achieves Oblivious Leader Selection assuming  $\delta \geq 2\alpha$ . Its fairness is analyzed below.

---

### Algorithm 2 Oblivious Leader Selection

---

**Input:**  $r$

- 1: Let  $r$  be given (representing the current iteration number in the outer protocol). Every honest participant  $P_i$  sends  $m_i = (i, \text{Sig}_i(r))$  to all other participants in its view.
- 2: Every honest participant  $P_i$  forwards messages with valid signatures to all other participants in the view of  $P_i$ .
- 3: Every honest participant  $P_i$  receives at most  $n$  forwarded messages from each  $P_j$  in its view (and ignore the messages after the first  $n$ ).  $P_i$  computes a set  $S_i$  of messages that are forwarded by at least  $(\delta - \alpha)n$  participants in its view, and send  $S_i$  to all participants in its view.
- 4: Every honest participant  $P_i$  receives a set  $S_j$  from every participant  $P_j$  in its view.
  - $P_i$  computes a set  $S_i^*$  of messages that appear in at least  $(1 - \alpha)n$  received sets.
  - For every  $m_k \in S_i^*$ ,  $P_i$  computes  $H(m_k)$  and outputs  $P_l$  where  $l$  is the smallest id such that for all  $m_k \in S_i^*$   $H(m_l) \leq H(m_k)$ .

**Output:**  $P_l$

---

We now show the following claims about Algorithm 2.

**Claim 6.** In any iteration  $r$ , if  $P_i, P_j$  are any two honest participants, then  $m_j = (j, \text{Sig}_j(r)) \in S_i^*$  in Step 4.

*Proof.* Consider any honest participant  $P_k$ . By our model assumption, there are at least  $(1 - \alpha)n$  honest participants in the view of  $P_j$ , and at least  $((1 - \alpha) - (1 - \delta))n = (\delta - \alpha)n$  of them are also in the view of  $P_k$ . Therefore, in **Step 3**  $P_k$  receives  $m_j$  forwarded by least  $(\delta - \alpha)n$  honest participants. Hence  $m_j \in S_k$ .

By the argument above, every honest  $P_k$  in the view of  $P_i$  sends  $S_k$  with  $m_j \in S_k$  in **Step 3**. Since there are at least  $(1 - \alpha)n$  of them, we have  $m_j \in S_i^*$  in **Step 4**.  $\square$

**Claim 7.** *In any iteration  $r$ , if  $P_i$  is honest, and if  $\delta > 2\alpha$ , then  $S_i$  contains at most  $2n$  messages from corrupted participants.*

*Proof.* We first calculate the total number of times any message from a corrupted participants gets forwarded to  $P_i$ . By our model assumption, there are at most  $\alpha n$  corrupted participants in the view of  $P_i$  that can each forward  $n$  corrupted messages. The rest  $(1 - \alpha)n$  honest participants will each forward at most  $\alpha n$  corrupted messages. In total, we get

$$\alpha n^2 + (1 - \alpha)\alpha n^2 = \alpha(2 - \alpha)n^2$$

For a corrupted message to be accepted into  $S_i$ , it must be forwarded at least  $(\delta - \alpha)n$  times. Therefore, the number of accepted corrupted messages is at most

$$\frac{\alpha(2 - \alpha)n^2}{(\delta - \alpha)n} \leq \frac{\alpha}{\delta - \alpha} 2n < 2n$$

□

The above claims lead to the following lemma that we will use as a building block to prove Theorem 1.

**Lemma 8.** *If  $\delta > 2\alpha$ , there exists a three round Oblivious Leader Selection protocol with fairness  $1/(5n)$ . Further if  $\alpha < 1/2 - \epsilon$  for some positive constant  $\epsilon$ , then there exists a three round Oblivious Leader Selection protocol with constant fairness.*

*Proof.* We show that Algorithm 2 is such a protocol. Let  $C$  be the set of all honest participants. Consider the union of all honest  $S_i^*$ :  $S^* = \cup_{P_i \in C} S_i^*$ . If  $P_l$  is an honest participant and  $l$  is the smallest id such that for all  $m_k \in S^*$   $H(m_l) \leq H(m_k)$ , then by Claim 6,  $P_l \in S_i^*$  for all honest  $P_i$ . Hence all honest  $P_i$  will output  $P_l$  in **Step 4** (i.e. an honest leader  $P_l$  is elected). Furthermore, by the definition of Random Oracle,  $H(m_k)$  are uniform random and independent for all  $m_k \in S^*$ . Therefore, the probability that an honest leader  $P_l$  is elected is exactly  $|C|/|S^*|$ . We now calculate the total number corrupted messages in  $S^*$ . In **Step 4**, any corrupted message ever accepted into an honest  $S_i^*$  must appears in at least  $(1 - \alpha)n$  sets, of which at least  $(1 - 2\alpha)n$  are from an honest participant. By Claim 7, the total number of corrupted messages in all honest  $S_i$  in **Step 3** is  $2nC$ . Therefore, the number of corrupted messages in  $S^*$  is at most

$$\frac{2n|C|}{(1 - 2\alpha)n} = \frac{2}{1 - 2\alpha}|C|$$

And we have  $|S^*| \leq |C| + \frac{2}{1 - 2\alpha}|C|$ . Since the number of corrupted participants in the view of any honest party is an integral number, we have  $1/2 - \alpha \geq 1/(2n)$ . The probability that an honest leader is elected is given by

$$\frac{|C|}{|S^*|} \geq \frac{|C|}{|C| + \frac{2}{1/(2n)}|C|} = \frac{1}{1 + 4n} \geq \frac{1}{5n}$$

In the case of  $\alpha < 1/2 - \epsilon$  for some constant  $\epsilon$ , we get

$$\frac{|C|}{|S^*|} \geq \frac{|C|}{|C| + \frac{2}{2\epsilon}|C|} = \frac{1}{1 + 1/\epsilon}$$

Which is a constant.

□

### 3.3 Main protocol

The main protocol is inspired by the one present in [Mic17]. At a high level the protocol ensures that if one honest participant decides to terminate with some output  $v \in \{0, 1\}$ , it is sure that no other honest participant terminates with a different  $v' \neq v$  in the same round, and that all honest participants will be able to terminate in the next round. When all honest participants hold the same value, they terminate immediately. When an honest leader is elected, they terminate with probability  $1/2$ . Overall, if an honest leader is elected with constant probability, then our main protocol terminates with expected constant round. Similarly, if an honest leader is elected with probability  $\Omega(1/n)$ , then our main protocol terminates with expected  $O(n)$  rounds. The main protocol is shown in Algorithm 3.

---

#### Algorithm 3 Byzantine Agreement

---

**Input:**  $v_i \in \{0, 1\}$ : the initial value;  $r \leftarrow 0$ : the current iteration;  $h_i \leftarrow 0$ : whether to halt.

- 1: Every honest  $P_i$  runs a Graded Broadcast protocol as the dealer with message  $v_i$ . In the end,  $P_i$  outputs  $(v_j, g_j)$  for every participant  $P_j$  in its view, and accepts only values with grade 1.
  - If  $h_i > 0$ , do nothing.
  - If more than  $(1 - 1/2\delta)n$  0s are accepted, then set  $v_i \leftarrow 0$  and  $h_i \leftarrow 1$ .
  - If more than  $(1 - 1/2\delta)n$  1s are accepted, then set  $v_i \leftarrow 1$ .
  - Otherwise, set  $v_i \leftarrow 0$ .
- 2: Every honest  $P_i$  runs a Graded Broadcast protocol as the dealer with message  $v_i$ . In the end,  $P_i$  outputs  $(v_j, g_j)$  for every participant  $P_j$  in its view, and accepts only values with grade 1.
  - If  $h_i > 0$ , do nothing.
  - If more than  $(1 - 1/2\delta)n$  1s are accepted, then set  $v_i \leftarrow 1$  and  $h_i \leftarrow 1$ .
  - If more than  $(1 - 1/2\delta)n$  0s are accepted, then set  $v_i \leftarrow 0$ .
  - Otherwise, set  $v_i \leftarrow 1$ .
- 3: Every honest  $P_i$  sends a random bit  $b \xleftarrow{R} \{0, 1\}$  to every participants in its view. In the end,  $P_i$  receives  $b_j$  from every participant  $P_j$  in its view.
- 4: Every honest  $P_i$  runs an Oblivious Leader Selection protocol with input  $r$  and outputs  $P_{l_i}$ .
- 5: Every honest  $P_i$  runs a Graded Broadcast protocol as the dealer with message  $v_i$ . In the end,  $P_i$  outputs  $(v_j, g_j)$  for every participant  $P_j$  in its view, and accepts only values with grade 1.
  - If  $h_i > 0$ , do nothing.
  - If more than  $(1 - 1/2\delta)n$  1s are accepted, then set  $v_i \leftarrow 1$ .
  - If more than  $(1 - 1/2\delta)n$  0s are accepted, then set  $v_i \leftarrow 0$ .
  - If  $P_{l_i}$  is in the view of  $P_i$ , then set  $v_i \leftarrow b_{l_i}$ .
- 6: Set  $r \leftarrow r + 1$ .
  - If  $h_i = 2$  halts with  $v^* = v_i$ .
  - If  $h_i = 1$ , sets  $h_i \leftarrow 2$ .

Go back to **Step 1**.

**Output:**  $v_i^* \in \{0, 1\}$

---

We now show the following claims about Algorithm 3.

**Claim 9.** (Validity) *If every honest participant  $P_i$  has the same initial value  $v$ , and if  $\delta > 2\alpha$ , then they all terminate in the second iteration.*

*Proof.* By our model assumption and the correctness of Graded Broadcast, every honest participant  $P_i$  in **Step 1** will output  $(v, 1)$  from at least  $(1 - \alpha)n$  honest Graded Broadcast. Since  $\delta > 2\alpha$ ,  $\vec{v}_i$  will contain more than  $(1 - 1/2\delta)n$  values of  $v$ .

If  $v = 0$ ,  $P_i$  sets  $h_i \leftarrow 1$  in **Step 1**. If  $v = 1$ ,  $P_i$  keeps  $v_i$  unchanged and sets  $h_i \leftarrow 1$  in **Step 2** by the same argument. Once  $h_i = 1$ ,  $v_i$  never changes and  $P_i$  halts in the next iteration with  $v^* = v_i$ .  $\square$

**Claim 10.** *If some honest participant  $P_i$  sets  $h_i \leftarrow 1$  in iteration  $r$  with  $v_i = v$ , then every other honest participant  $P_j$  will have set  $h_j \leftarrow 1$  with  $v_j = v$  by the end of iteration  $r + 1$ , and all of them halts by the end of iteration  $r + 2$ .*

*Proof.* It suffice to consider the first honest  $P_i$  that sets  $h_i \leftarrow 1$ .

- Suppose this happened in **Step 1**, consider a different honest participant  $P_j$ , who didn't set  $h_j \leftarrow 1$  in **Step 1** (otherwise, we are done). By our model assumption, of any  $(1 - 1/2\delta)n$  participants who Graded Broadcasted messages 0 to  $P_i$ , at least  $((1 - 1/2\delta) - (1 - \delta))n = 1/2\delta n$  of them are also in the view of  $P_j$ . By correctness of Graded Broadcast,  $P_j$  cannot accept more than  $(1 - 1/2\delta)n$  values of 1 in **Step 1**. Hence  $P_j$  sets  $v_j \leftarrow 0$ . That is, by the end of **Step 1**, every honest party  $P_j$  holds  $v_j = 0$ . Repeating the argument from Claim 9,  $P_j$  keeps  $v_j$  unchanged in this iteration. If  $P_j$  didn't set  $h_j \leftarrow 1$  in **Step 1**, it will in the next iteration.
- Suppose this happened in **Step 2**. By assumption,  $P_i$  is the first honest party who sets  $h_i \leftarrow 1$ , and no other honest  $P_j$  have set  $h_j \leftarrow 1$  in **Step 1**. The rest follows from exactly the argument in the previous case.

$\square$

Note that by Claim 10, honest participants always halt with the same value. It now suffice to only consider the case when no honest participant  $P_i$  has set  $h_i \leftarrow 1$ :

**Claim 11.** *Suppose no honest participant has set  $h_i \leftarrow 1$  in some iteration  $r$ . If an honest leader  $P_l$  is elected in iteration  $r$ , and if  $\delta > 2\alpha$  then with probability  $1/2$   $P_l$  will set  $P_i$  in iteration  $r + 1$ .*

*Proof.* Suppose some honest  $P_i$  in the view of  $P_l$  in **Step 5** accepted more than  $(1 - 1/2\delta)n$  values of some value  $v$ . By the argument from Claim 10, no other honest  $P_j$  have accepted more than  $(1 - 1/2\delta)n$  values of the different value  $v^c$ . That is, in **Step 5** the honest participants in the view of  $P_l$  either set their values to  $v$ , or to  $b_l$ , the random bit from  $P_l$ .

With probability  $1/2$ ,  $b_l = v$ , and all honest participants in the view of  $P_l$  hold the same value in the next iteration. Since there are at least  $(1 - \alpha)n$  of them, and  $\delta > 2\alpha$ ,  $P_l$  will accept more than  $(1 - 1/2\delta)n$  values of  $v$  and sets in  $h_l \leftarrow 1$  in the next iteration.  $\square$

Theorem 1 now follows directly from Lemma 5, Lemma 8, Claim 9, Claim 10, and Claim 11.

# Chapter 4

## Negative Result

### 4.1 Broadcast from Byzantine Agreement

### 4.2 Negative Result for Broadcast

**The case of  $\alpha \geq 1/2$**  We start with the case where  $\alpha \geq 1/2$ . Let  $p, r$  be positive integers. Without loss of generality, assume  $p > r$ . Figure shows a configuration where a total of  $N = r + 2p$  participants are divided into 4 groups,  $A, B, C$  and  $F$ .

**The case of  $\delta \leq 2\alpha$**  asdf

November 20, 2019  
DRAFT

## **Chapter 5**

### **Extension: Unidirectional Edges**



November 20, 2019  
DRAFT

# Bibliography

- [BO83] Michael Ben-Or. Another Advantage of Free Choice (Extended Abstract): Completely Asynchronous Agreement Protocols. In *Proceedings of the Second Annual ACM Symposium on Principles of Distributed Computing*, PODC '83, pages 27–30, New York, NY, USA, 1983. ACM. event-place: Montreal, Quebec, Canada. 1.1
- [CGO10] Nishanth Chandran, Juan Garay, and Rafail Ostrovsky. Improved Fault Tolerance and Secure Computation on Sparse Networks. In David Hutchison, Takeo Kanade, Josef Kittler, Jon M. Kleinberg, Friedemann Mattern, John C. Mitchell, Moni Naor, Oscar Nierstrasz, C. Pandu Rangan, Bernhard Steffen, Madhu Sudan, Demetri Terzopoulos, Doug Tygar, Moshe Y. Vardi, Gerhard Weikum, Samson Abramsky, Cyril Gavoille, Claude Kirchner, Friedhelm Meyer auf der Heide, and Paul G. Spirakis, editors, *Automata, Languages and Programming*, volume 6199, pages 249–260. Springer Berlin Heidelberg, Berlin, Heidelberg, 2010. 1.1
- [Dol82] Danny Dolev. The Byzantine generals strike again. *Journal of Algorithms*, 3(1):14–30, March 1982. 1.1
- [DS83] D. Dolev and H. R. Strong. Authenticated Algorithms for Byzantine Agreement. *SIAM Journal on Computing*, 12(4):656–666, November 1983. 1.1
- [FL82] Michael J. Fischer and Nancy A. Lynch. A lower bound for the time to assure interactive consistency. *Information Processing Letters*, 14(4):183–186, June 1982. 1.1
- [FM97] Pesech Feldman and Silvio Micali. An optimal probabilistic protocol for synchronous byzantine agreement. *SIAM Journal on Computing*, 26(4):873–933, 1997. 1.1, 3.1
- [GM98] Juan A. Garay and Yoram Moses. Fully Polynomial Byzantine Agreement for Processors in Rounds. *SIAM J. Comput.*, 27(1):247–290, February 1998. 1.1
- [GO92] Shafi Goldwasser and Rafail Ostrovsky. Invariant signatures and non-interactive zero-knowledge proofs are equivalent. In *Annual International Cryptology Conference*, pages 228–245. Springer, 1992. 2.2
- [KK06] Jonathan Katz and Chiu-Yuen Koo. On expected constant-round protocols for byzantine agreement. In *Advances in Cryptology - CRYPTO*, Lecture Notes in Computer Science, pages 445–462. Springer Berlin Heidelberg, 2006. 1.1, 2.4, 3.2
- [KS09] Valerie King and Jared Saia. From almost everywhere to everywhere: Byzantine agreement with  $\tilde{O}(n^{3/2})$  bits. In *Distributed Computing*, Lecture Notes in Computer Science, pages 464–478. Springer Berlin Heidelberg, 2009. 1.1

- [LSP] Leslie Lamport, Robert Shostak, and Marshall Pease. The Byzantine Generals Problem. *ACM Transactions on Programming Languages and Systems*, 4(3):20. 1.1, 1.1
- [Mic17] Silvio Micali. Byzantine agreement, made trivial. 2017. 1.1, 2.4, 3.2, 3.3
- [MRVil] S. Micali, M. Rabin, and S. Vadhan. Verifiable random functions. In *40th Annual Symposium on Foundations of Computer Science (Cat. No.99CB37039)*, page nil, -nil. 2.2, 2.3
- [PSL80] M. Pease, R. Shostak, and L. Lamport. Reaching agreement in the presence of faults. *Journal of the ACM*, 27(2):228–234, April 1980. 1.1, 1.1
- [Rab83] Michael O. Rabin. Randomized byzantine generals. In *24th Annual Symposium on Foundations of Computer Science (sfcs 1983)*, pages 403–409, November 1983. ISSN: 0272-5428. 1.1